



# CMP636 Distributed Systems


## Static Sharded KV Server Lab

Ayman AboElHassan, PhD  
Assistant Professor

[ayman.abo.elmaaty@eng.cu.edu.eg](mailto:ayman.abo.elmaaty@eng.cu.edu.eg)



# Basic Idea





# Basic Idea

1. Client send request to KV server
2. Hash function determines which group holds the key-value pair
3. KV server **shard** send operation to Raft
4. Raft reaches value consensus → Total Broadcast to followers
5. Followers reply to leader
6. Once quorum Ack received → Commit KV change





# Requirement

Change last lab key/value server

1. Create a shard manager server
2. Create **2 sets of “3 servers”**
3. Each server maintains its own in-memory map of key/value pairs (Local KV Map)
4. Each **group** maintain a set of different key/value pairs (**sharding**)
5. Clients sends read/write/append to 1 of the  $3 \times 2$  servers (according to the sharded group)
6. Servers of same group reach consensus on KV operation through Raft



# Shard Manager Server

A logical server which knows which group of servers has which key

- Basic implementation “Static”  $\rightarrow$  Hash function map  $\rightarrow \% \text{ nGroups}$

## FNV-1 hash [\[edit\]](#)

The FNV-1 hash algorithm is as follows:[\[9\]](#)[\[10\]](#)

```
algorithm fnv-1 is
    hash := FNV_offset_basis

    for each byte_of_data to be hashed do
        hash := hash × FNV_prime
        hash := hash XOR byte_of_data

    return hash
```





# Python Implementation

## Server

1. Create SyncObj
2. Create empty Replicated Dict
3. Create 1 RPC server stub
4. Wait for client requests
5. Perform client's request and print the operation input/result

## Client

1. Create 5 client threads
2. Each thread
  1. Select a random request out of:
    - **Get** value of K@
    - **Put** value ClientNum in K@
    - **Append** value ClientNum to K@
  2. **Get server group from shard manager**
  3. Send a request to the target group
  4. Wait for response





# Python Implementation

## gRPC proto “kv server”

- Service1: **Get**
  - Message: **key**
  - Behavior:  
return map[key]
  - Response: **value**
- Service2: **Put**
  - Message: **key, value**
  - Behavior:  
map[key] = value
  - Response: **success**
- Service3: **Append**
  - Message: **key, args**
  - Behavior:  
old\_value = map[key]  
map[key] += args
  - Response: **old\_value**





# Python Implementation

gRPC proto “**Shard Manager**”

- Service 1: **GetShardIndex**
- Message: **key**
- Behavior:  
return Shard's Group gRPC port ID
- Response: **gRPC port ID**







# Why Kubernetes?

1. Dynamic deployment “port assignment”
2. Dynamic scale-out “vertical scale-up”
3. Restart on failure
4. Roll-out updates





Thank you

